

CS 398 ACC

Spark SQL

Prof. Robert J. Brunner

Ben Congdon
Tyler Kim

What's going on with the cluster?

People running “local” jobs on Master consumes disproportionate amount of CPU

- If master is unresponsive, it makes the entire cluster useless
- Please be courteous of other students during “peak” hours
 - We will be more aggressive in kicking out jobs if the problem continues

Course Cluster

Back-up / secondary cluster will be available.

Check the Cluster page on the website

- Same SSH key

Outline

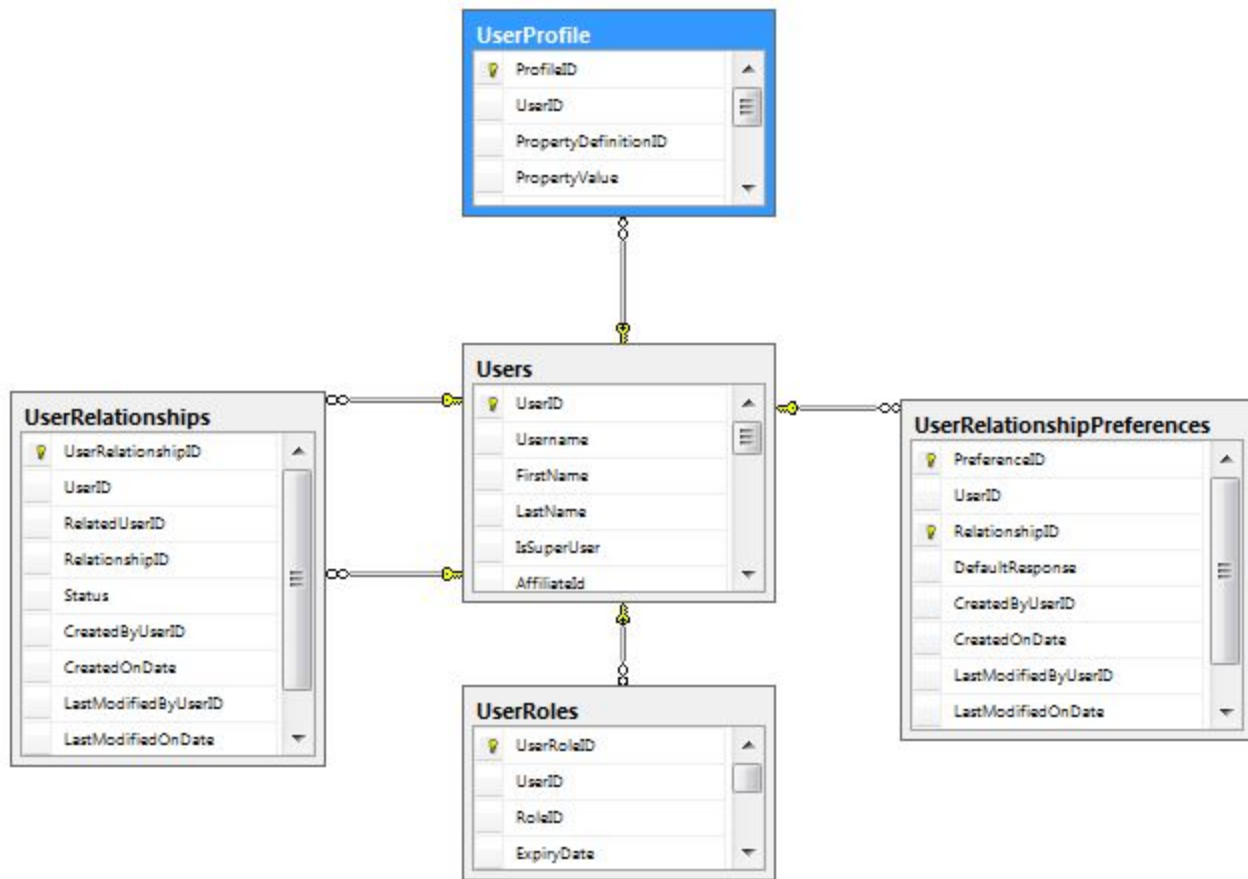
- Traditional Databases
- SQL
 - Optimizations
- Spark SQL

Outline

- **Traditional Databases**
- SQL
 - Optimizations
- Spark SQL


RDBMS

- Relational Database Management Systems
 - Systems that deal with relational data (data that points to other data)
- A database management system manages how the data is stored and retrieved. Usually the data is modified with SQL
- E.g: MySQL, PostgreSQL, OracleDB, etc



Other Features

- RDBMS handles data backups, logically storing data, distributing data to leader followers, permissions, data integrity, handling and load balancing queries, and optimization.
- RDBMSs do all of this “under the hood” (mostly)



RDBMS Types of Data


- RDBMSs like simple data: INTEGERS, STRINGS, etc
 - They don't like handling JSON, HASHMAP, LISTS
 - Complex data types are more difficult for the SQL engine to optimize against
- If you think you need advanced data type functionality:
 - **Seriously rethink your application design**
- If you are absolutely sure that you need it:
 - You should probably use another application server.

Outline

- Traditional Databases
- **SQL**
 - Optimizations
 - Spark SQL

Structured Query Language

- Most of you have had some interaction with SQL
- SQL was made for both programmers and for accountants who were used to spreadsheets
- We can imagine taking data from spreadsheets, join from different sheets etc



Basic Commands - Data Definition Language (DDL)

- DDL lets you create, destroy, alter, and modify constraints on data
- You can think of them as operations that set up where data will go

- ```
CREATE TABLE (
 id INTEGER, name VARCHAR(255), location VARCHAR(255)
);
```
- ```
ALTER TABLE ADD status INTEGER;
```
- ```
ALTER TABLE ADD blah INTEGER NOT NULL;
```
- ```
DROP TABLE;
```


Data Modification Language - DML

- This adds, deletes, selects, and updates data (basic CRUD operations)
 - This lets you put data into the database tables
-
- `INSERT INTO table (col1, col2, ..) VALUES (v1, v2, ..), ..`
 - `DELETE FROM table where col1 = ...`
 - `UPDATE table SET col1='asdf' WHERE col2='asd'`
 - `SELECT * FROM table`


Data Modification Language Extensions

- The data modification language also lets you do more powerful things when retrieving data
 - We can have data GROUP BY a certain column(s)
 - Have data ORDER BY some column(s)
 - We can JOIN multiple spreadsheets based on a column
- We can have SQL calculate functions or aggregations on the fly
- Usually RDBMSs are optimized for read-heavy workloads


CHEATSHEET SQL JOINS




```
SELECT <auswahl>
FROM tabelleA A
LEFT JOIN tabelleB B
ON A.key = B.key
```




```
SELECT <auswahl>
FROM tabelleA A
INNER JOIN tabelleB B
ON A.key = B.key
```




```
SELECT <auswahl>
FROM tabelleA A
RIGHT JOIN tabelleB B
ON A.key = B.key
```




```
SELECT <auswahl>
FROM tabelleA A
LEFT JOIN tabelleB B
ON A.key = B.key
WHERE B.key IS NULL
```



```
SELECT <auswahl>
FROM tabelleA A
RIGHT JOIN tabelleB B
ON A.key = B.key
WHERE A.key IS NULL
```



```
SELECT <auswahl>
FROM tabelleA A
FULL OUTER JOIN tabelleB B
ON A.key = B.key
```



```
SELECT <auswahl>
FROM tabelleA A
FULL OUTER JOIN tabelleB B
ON A.key = B.key
WHERE A.key IS NULL
OR B.key IS NULL
```

SQL Prepared Statements

- Actual interactions with the database.
 - `INSERT INTO table VALUES (`+userid+`);`
- What if `userid = "1; SELECT * FROM table WHERE col1 NOT IN ("?`
 - `INSERT INTO table VALUES (1); SELECT * FROM table WHERE col1 NOT IN ();`
 - This will give us back all the results from the database!

SQL Prepared Statements

- To avoid this, we have prepared statements
- `INSERT INTO table VALUES (?)` and, send userid separately
- This avoids the injection problem but doesn't let SQL server optimize database queries

Outline

- Traditional Databases
- SQL
 - **Optimizations**
- Spark SQL

SQL Turing Completeness

- Every SQL statement (in ANSI SQL) will terminate
- The Non-Turing Completeness of SQL let's us optimize many portions of queries

User tips for optimizing SQL queries

- Don't use ``SELECT *`` statements, you usually are selecting more rows than need be
- If you have multiple levels of joins then you may want to consider staging your data into an intermediate table in order to reduce communication overhead
- Add indices! Indices can slow updates but drastically speed up complex queries if the indices are on the appropriate columns

SQL Optimizer: Prediction

- Consider a query like ``select col1 from table where col1=1 AND col2=2;``
- Your server has the choice of filtering by col2 and then col1 or by col1 then col2.
- If the server knows that there are a lot of NULL values in col2 which would reduce the number of rows in consideration a lot, it will filter based on col2 first and then filter on col1 because the complexity will be $\text{NUM_ROWS} * \text{SMALL_NUMBER}$

SQL Optimizer: Lazy Joins

- A join is when you combine two tables on a column

c1	c2
1	2
2	4

c3	c4
1	1
2	1
3	2

Example Join

```
SELECT * FROM t1 JOIN t2 USING (c1, c4);
```

c1	c2	c3	c4
1	2	1	1
1	2	2	1
2	4	3	2

Lazy Join

- SQL may filter the data before joining, may group by before joining if you know that one of the columns is in one of the table
- This is very ad-hoc prediction because SQL usually doesn't keep track of super in depth statistics
- As a SQL server runs longer, then it gets better at this prediction
 - The main reason that it can't keep track of all of this information is due to concurrency bottlenecks so it makes static analyses instead

Outline

- Traditional Databases
- SQL
 - Optimizations
- **Spark SQL**

Spark SQL


- Distributed in-memory computation on massive scale (Just like Spark!)
- Can use all data sources that Spark supports natively:
 - Can import data from RDDs
 - JSON/CSV files can be loaded with inferred schema
 - Parquet files - Column-based storage format
 - Supported by many Apache systems (big surprise!)
 - Hive Table import
 - A popular data warehousing platform by Apache

Spark SQL

- SQL using Spark as a “Database”
 - Spark SQL is best optimized for retrieving data
 - Don't UPDATE, INSERT, or DELETE
- Optimization handled by a newer optimization engine, **Catalyst**
 - Creates physical execution plan and compiles directly to JVM bytecode
- Can function as a compatibility layer for firms that use RDBMS systems

Spark DataFrames



- Dataset organized into named columns
- Similar to structure as Dataframes in Python (i.e. Pandas) or R
- Lazily evaluated like normal RDDs
- Tends to be more performant than raw RDD operations



Pandas DataFrame

pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



Does in-memory computing, but:

- Not scalable by itself.
- Not fault tolerant.

```
import pandas as pd
```

```
df = pd.read_csv("/path/to/data.json")
```

```
df
```

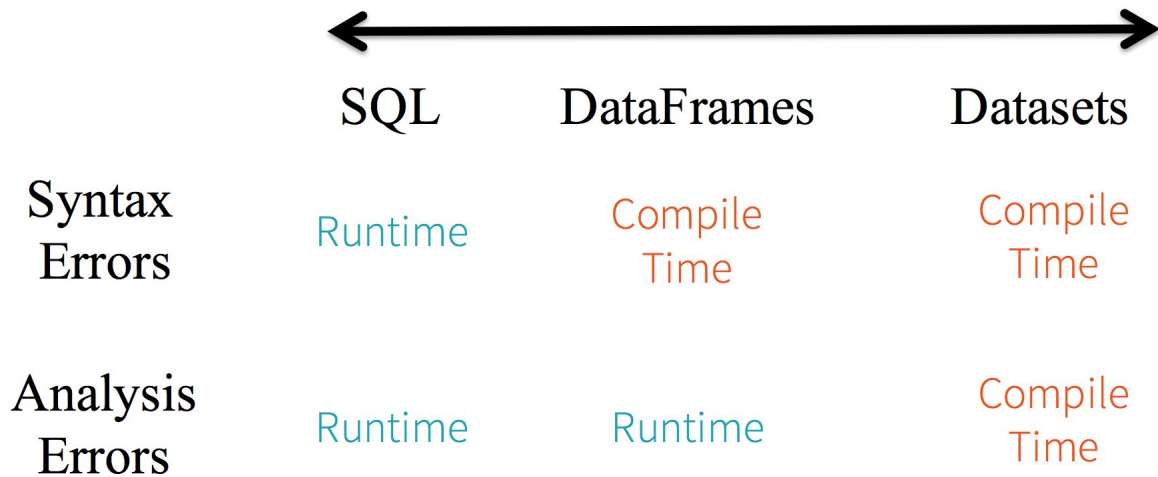
	first_name	last_name	age	preTestScore	postTestScore
0	Jason	Miller	42	4	25,000
1	Molly	Jacobson	52	24	94,000
2	Tina	.	36	31	57
3	Jake	Milner	24	.	62
4	Amy	Cooze	73	.	70

Spark DataFrames

- When to prefer RDDs over DataFrames:
 - Need low-level access to data
 - Data is mostly unstructured or schemaless
- When to prefer DataFrames over RDDs:
 - Operations on structured data
 - If higher-level abstractions are useful (i.e. joins, aggregation, etc.)
 - High-performance is desired, and workload fits within DataFrame APIs
 - Catalyst optimization makes DataFrames more performant on average

Spark DataSets

- Strongly-typed DataFrames
- Only accessible in Spark2+ using Scala
- Operations on DataFrames are all statically typed, so you catch type errors at compile-time



Data Ingest (RDD)

```
from pyspark.sql import SQLContext

sqlContext = SQLContext(sc)

users_rdd = sc.parallelize([[1, 'Alice', 10], [2, 'Bob', 8]])

users = sqlContext.createDataFrame(
    users_rdd,
    ['id', 'name', 'num_posts'])

users.printSchema()

#root
# |-- id: long (nullable = true)
# |-- name: string (nullable = true)
# |-- num_posts: long (nullable = true)
```


Data Ingest (JSON)

```
from pyspark.sql import SQLContext
sqlContext = SQLContext(sc)

users = sqlContext.read.json("/path/to/users.json")

users.printSchema()

# root
#   |-- id: long (nullable = true)
#   |-- name: string (nullable = true)
#   |-- num_posts: long (nullable = true)
```

SQL API

```
# Register users DataFrame as a table called "users"
users.createOrReplaceTempView( 'users' )
```

```
# Query the table
sqlContext.sql(
    'SELECT * FROM users WHERE name="Bob"'
).collect()
```

```
# [Row(id=2, name='Bob', num_posts=8)]
```

DataFrame API

```
# Same query can be done with DataFrame API
```

```
users.filter(users.name=='Bob').collect()
```

```
# [Row(id=2, name='Bob', num_posts=8)]
```

```
users.filter(users.name=='Eve').select('num_posts').collect()
```

```
# [Row(num_posts=10)]
```